

Cilk Plus Reducers

Albert DeFusco

April 7, 2015

Parallel Gold Sifting

- ▶ Each pan can sift a constant amount of dirt/day
- ▶ More pans means more dirt sifted

Parallel Gold Sifting

- ▶ Each pan can sift a constant amount of dirt/day
- ▶ More pans means more dirt sifted
- ▶ Each pan sifts independently
- ▶ Each pan has a defined amount of dirt to sift

Parallel Gold Sifting

- ▶ Each pan can sift a constant amount of dirt/day
- ▶ More pans means more dirt sifted
- ▶ Each pan sifts independently
- ▶ Each pan has a defined amount of dirt to sift
- ▶ sifting is parallelizable

Serial Gold Sifting

```
1  #include <list>
2  class pan
3  {
4      public:
5          pan();           //create array of random integers between 1 and 10,000
6          int sift();       //returns frequency of the number 79 (atomic number of gold)
7          bool hasGold();  //calls sift; returns true if sift() > 0
8  }
9
10 int main()
11 {
12     std::list<int> withGold;
13     Pan* myPans = new Pan[nPans];
14
15     for(int i=0; i<nPans; ++i)
16     {
17         bool gold = myPans[i].hasGold();
18         if(gold) {
19             withGold.push_back(i);
20         }
21     }
22     std::list<int>::const_iterator iterator;
23     for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
24         std::cout << *iterator << " ";
25     std::cout << endl;
26
27     return 0;
28 }
```

Gold Rush

Gold Rush!

10000 total chunks of dirt

1000 pans

Found gold in 15 pans

Pan IDs: 94 142 265 268 289 440 442 443 569 600 721 781 783 806 818

serial execution took 5.60495 seconds

Parallel Gold Sifting

```
1  #include <list>
2  #include <cilk/cilk.h>
3  class pan
4  {
5      public:
6          pan();           //create array of random integers between 1 and 10,000
7          int sift();       //returns frequency of the number 79 (atomic number of gold)
8          bool hasGold();  //calls sift; returns true if sift() > 0
9  }
10
11 int main()
12 {
13     std::list<int> withGold;
14     Pan* myPans = new Pan[nPans];
15
16     cilk_for(int i=0; i<nPans; ++i)
17     {
18         bool gold = myPans[i].hasGold();
19         if(gold) {
20             withGold.push_back(i);
21         }
22     }
23     std::list<int>::const_iterator iterator;
24     for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
25         std::cout << *iterator << " ";
26     std::cout << endl;
27
28     return 0;
29 }
```

Parallel Gold Sifting

- ▶ There is a problem to be solved
 - ▶ How do we keep track of which pans have gold?
- ▶ Where does the result get stored?
- ▶ Who can access the result?

Parallel Gold Sifting

```
1  #include <list>
2  #include <cilk/cilk.h>
3  class pan
4  {
5      public:
6          pan();           //create array of random integers between 1 and 10,000
7          int sift();       //returns frequency of the number 79 (atomic number of gold)
8          bool hasGold();  //calls sift; returns true if sift() > 0
9  }
10
11 int main()
12 {
13     std::list<int> withGold;
14     Pan* myPans = new Pan[nPans];
15
16     cilk_for(int i=0; i<nPans; ++i)
17     {
18         bool gold = myPans[i].hasGold();
19         if(gold) {
20             withGold.push_back(i);
21         }
22     }
23
24     std::list<int>::const_iterator iterator;
25     for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
26         std::cout << *iterator << " ";
27     std::cout << endl;
28
29     return 0;
30 }
```

may not be thread-safe

Thread Safety

- ▶ Unsafe operations
 - ▶ Multiple threads accessing the same address
 - ▶ Basic types are not thread safe
 - ▶ STL containers *may* be thread safe for some operations

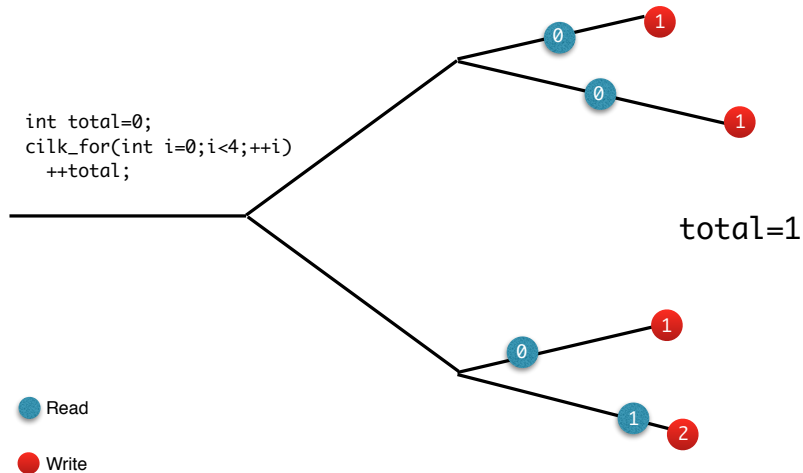
Thread Safety

- ▶ Unsafe operations
 - ▶ Multiple threads accessing the same address
 - ▶ Basic types are not thread safe
 - ▶ STL containers *may* be thread safe for some operations
- ▶ Threads read and write memory at undetermined times

Thread Safety

- ▶ Unsafe operations
 - ▶ Multiple threads accessing the same address
 - ▶ Basic types are not thread safe
 - ▶ STL containers *may* be thread safe for some operations
- ▶ Threads read and write memory at undetermined times
- ▶ Leads to a race condition

Race Condition



Inefficient solutions

- ▶ Locking
 - ▶ Requires careful programming
 - ▶ Non deterministic
- ▶ cannot use `cilk_sync` in the loop
 - ▶ will only sync child threads, not all threads
- ▶ Break the loop
 - ▶ Requires more storage and management

```
1  #include <cilk / cilk.h>
2
3  double *sum = new double[N];
4  // parallel
5  cilk_for (int i=0; i<N; ++i)
6      sum[N] = f(N);
7
8  // serial
9  int total=0;
10 for (int i=0; i<N; ++i)
11     total+=sum[N];
```

Cilk Reducers

- ▶ Provide thread safe access to a “smart pointer”
- ▶ Any associative operation is a valid reducer

$$x \text{ OP } y = y \text{ OP } x$$

- ▶ Small performance overhead for usage
- ▶ Very extensible in C++
- ▶ Operations are guaranteed to execute in the same order as in serial

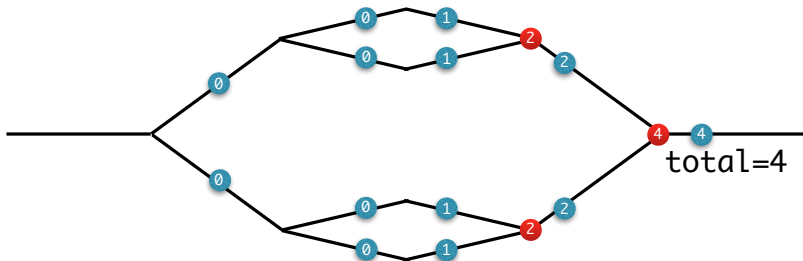
Cilk Reducers: views

```
1 #include <cilk/reducer_opadd.h>
2 int total=0;
3 cilk::reducer<cilk::<op_add<int>> reducer_total (0);
4 for (int i=0; i<4; ++i)
5     ++*reducer_total;
6 total = reducer_total.get_value();
```

- ▶ At spawn each strand gets a private *view* of the reducer
- ▶ When strands merge
 - ▶ *views* are combined by *OP*
 - ▶ The combined *view* is given to the exit strand

Cilk Reducers: views

```
#include <cilk/reducer_opadd.h>
int total=0;
cilk::reducer<cilk::<op_add<int>> reducer_total (0);
for(int i=0;i<4;++i)
    ++*reducer_total;
total = reducer_total.get_value();
```



● Private View

● Merge update

Cilk Reducers: types

usage: `cilk::reducer<cilk::REDUCER_TYPE<<my_type>> my_reducer;`

Reducer Type	Description	Header
<code>op_add</code>	<code>++</code> , <code>--</code> , <code>+=</code> , <code>-=</code> , <code>+</code> , <code>-</code>	<code>#include <cilk/reducer_add.h></code>
<code>op_vector</code>	Provides <code>push_back()</code>	<code>#include <cilk/reducer_vector.h></code>
<code>op_list_append</code>	Provides <code>push_back()</code>	<code>#include <cilk/reducer_list.h></code>
<code>op_list_prepend</code>	Provides <code>push_front()</code>	<code>#include <cilk/reducer_list.h></code>
<code>op_max</code>	Returns maximum	<code>#include <cilk/reducer_max.h></code>
<code>op_min</code>	Returns minimum	<code>#include <cilk/reducer_min.h></code>
<code>op_ostream</code>	Provides <code><<</code>	<code>#include <cilk/reducer_string.h></code>

Table: <https://software.intel.com/en-us/node/522606>

Parallel gold sifting

```
1  #include <cilk/cilk.h>
2  #include <cilk/reducer_list.h>
3
4  std::list<int> withGold;
5  cilk::reducer< cilk::op_list_append<int> > reducer_withGold;
6  Pan* myPans = new Pan[nPans];
7
8  cilk_for(int i=0; i<nPans; ++i)
9  {
10     bool gold = myPans[i].hasGold();
11     if(gold) {
12         reducer_withGold->push_back(i);
13     }
14 }
15 withGold = reducer_withGold.get_value();
16
17 std::list<int>::const_iterator iterator;
18 for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
19     std::cout << *iterator << " ";
20 std::cout << endl;
```

Gold Rush

```
$>cat /proc/cpuinfo | grep Xeon | uniq -c
      16 model name      : Intel(R) Xeon(R) CPU E5-2670 0 @ 2.60GHz
$>CILK_WORKERS=16 ./goldRush
Gold Rush!
```

```
100000 total chunks of dirt
1000 pans
```

```
Found gold in 15 pans
```

```
Pan IDs: 94 142 265 268 289 440 442 443 569 600 721 781 783 806 818
```

```
Cilk identified the correct pans
```

```
serial execution took 5.60154 seconds
```

```
parallel execution took 0.39801 seconds with 16 workers
```

```
parallel speedup 14.0739
```

```
parallel efficiency 0.879616
```