

Cilk Plus Reducers

Albert DeFusco

April 6, 2015

Parallel Gold Sifting

- ▶ each pan can sift a constant amount of dirt/day
- ▶ more pans means more dirt sifted
- ▶ each pan sifts independently
- ▶ each pan has a definite amount of dirt to sift
- ▶ sifting is parallelizable

Serial Gold Sifting

```
1  #include <list>
2  class pan
3  {
4      public:
5          pan();           //create array of random integers between 1 and 10,000
6          int sift();       //returns frequency of the number 79 (atomic number of gold)
7          bool hasGold();  //calls sift; true if sift returns >0
8  }
9
10 int main()
11 {
12     std::list<int> withGold;
13     pan* myPans = new pan[nPans];
14
15     for(int i=0; i<nPans; ++i)
16     {
17         bool gold = myPans[i].hasGold();
18         if(gold) {
19             withGold.push_back(i);
20         }
21     }
22     std::list<int>::const_iterator iterator;
23     for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
24         std::cout << *iterator << " ";
25     std::cout << endl;
26
27     return 0;
28 }
```

Gold Rush

Gold Rush!

10000 total chunks of dirt

1000 pans

Found gold in 15 pans

Pan IDs: 94 142 265 268 289 440 442 443 569 600 721 781 783 806 818

serial execution took 5.60495 seconds

Parallel Gold Sifting

```
1  #include <list>
2  class pan
3  {
4  public:
5      pan();           //create array of random integers between 1 and 10,000
6      int sift();       //returns frequency of the number 79 (atomic number of gold)
7      bool hasGold();  //calls sift; true if sift returns >0
8  }
9
10 int main()
11 {
12     std::list<int> withGold;
13     pan* myPans = new pan[nPans];
14
15     cilk_for(int i=0; i<nPans; ++i)
16     {
17         bool gold = myPans[i].hasGold();
18         if(gold) {
19             withGold.push_back(i);
20         }
21     }
22     std::list<int>::const_iterator iterator;
23     for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
24         std::cout << *iterator << " ";
25     std::cout << endl;
26
27     return 0;
28 }
```

Parallel Gold Sifting

- ▶ Need to report which pans have gold
- ▶ How do we keep track of which pans have gold?


Parallel Gold Sifting

```
1  #include <list>
2  class pan
3  {
4  public:
5      pan();           //create array of random integers between 1 and 10,000
6      int sift();       //returns frequency of the number 79 (atomic number of gold)
7      bool hasGold();  //calls sift; true if sift returns >0
8  }
9
10 int main()
11 {
12     std::list<int> withGold;
13     pan* myPans = new pan[nPans];
14
15     cilk_for(int i=0; i<nPans; ++i)
16     {
17         bool gold = myPans[i].hasGold();
18         if(gold) {
19             withGold.push_back(i);
20         }
21     }
22
23     std::list<int>::const_iterator iterator;
24     for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
25         std::cout << *iterator << " ";
26     std::cout << endl;
27
28     return 0;
29 }
```

Thread Safety

- ▶ Unsafe operations
 - ▶ Multiple threads accessing the same address
 - ▶ Basic types are not thread safe
 - ▶ STL types are not thread safe
 - ▶ Threads read and write memory at undetermined times
 - ▶ Leads to a race condition

Race Condition



Inefficient solutions

- ▶ Non deterministic
 - ▶ Lock access with mutex
 - ▶ Requires careful programming
- ▶ cannot use `cilk_sync` in the loop
 - ▶ will only sync child threads, not all threads
- ▶ Break the loop
 - ▶ Requires more storage and management

```
1  #include <cilk / cilk.h>
2
3  double *sum = new double[N];
4  cilk_for (int i=0; i<N; ++i)
5      sum[N] = f(N)
6  double total=0.0;
7  for (int i=0; i<N; ++i)
8      total+=sum[N];
```

Cilk Reducers

- ▶ Provide thread safe access to a “smart pointer”
- ▶ Any associative operation is a valid reducer

$$x \text{ OP } y = y \text{ OP } x$$


- ▶ Small parallel overhead for usage
- ▶ Very extensible in C++
- ▶ Operations are guaranteed to execute in the same order as in serial

Cilk Reducers: views

- ▶ At spawn each strand gets a private “view” of the reducer
- ▶ When strands merge
 - ▶ “views” are combined by *OP*
 - ▶ The combined “view” is given to the exit strand

Cilk Reducers: views

```
int total=0;  
cilk::reducer<cilk::<op_add<int>> reducer_total (0);  
for(int i=0;i<4;++i)  
    ++*reducer_total;  
total = reducer_total.get_value();
```



● Private View

● Merge update

Parallel gold sifting

```
1  #include <cilk/cilk.h>
2  #include <cilk/reducer_list.h>
3
4  std::list<int> withGold;
5  cilk::reducer< cilk::op_list_append<int> > reducer_withGold;
6  pan* myPans = new pan[nPans];
7
8  cilk_for(int i=0; i<nPans; ++i)
9  {
10     bool gold = myPans[i].hasGold();
11     if(gold) {
12         reducer_withGold->push_back(i);
13     }
14 }
15 withGold = reducer_withGold.get_value();
16
17 list<int>::const_iterator iterator;
18 for (iterator = withGold.begin(); iterator != withGold.end(); ++iterator)
19     cout << *iterator << " ";
20 cout << endl;
```

Gold Rush

```
$>cat /proc/cpuinfo | grep Xeon | uniq -c
      16 model name      : Intel(R) Xeon(R) CPU E5-2670 0 @ 2.60GHz
$>CILK_WORKERS=16 ./goldRush
Gold Rush!
```

```
100000 total chunks of dirt
1000 pans
```

```
Found gold in 15 pans
```

```
Pan IDs: 94 142 265 268 289 440 442 443 569 600 721 781 783 806 818
```

```
Cilk identified the correct pans
```

```
serial execution took 5.60154 seconds
```

```
parallel execution took 0.39801 seconds with 16 workers
```

```
parallel speedup 14.0739
```

```
parallel efficiency 0.879616
```